

# PRÁCTICA Nº 4.4

## Agente para Registrar Cambios en un Fichero de Log



- Nombre y apellidos: Alvaro Lucio-Villegas de Cea



## Índice

---

Preparar el Fichero de Configuración.	3
Preguntas:	5
Prepara el fichero log.txt:	6
Crea el Agente Flume:	6
Modifica el Fichero /temp/log.txt	8

Vamos a crear un agente Flume muy simple, de tal forma que registre eventos que vamos a generar modificando un fichero de texto, como si fuera un fichero de Log. Para ello vamos a crear un archivo llamado log.txt en el directorio /tmp. El agente mostrará por la consola los cambios realizados en el archivo con el comando tail.

### Preparar el Fichero de Configuración.

- Crea un directorio flumeconf en la carpeta /home/cloudera.
- Dentro de ese directorio, crea el fichero agentefichero.log.conf.
- Utiliza como nombre del agente agentefichero.log
- Ese fichero debe contener todos los elementos de configuración que hemos visto en teoría. Vamos a concretarlos:

```
[cloudera@quickstart ~]$ ls
cloudera-manager  Downloads  kerberos  Public
cm_api.py         eclipse    lib        telefonica.java
copiaempleados.java empleados.java Music      Templates
Desktop           enterprise-deployment.json parcels     Videos
Documents         express-deployment.json Pictures    workspace
[cloudera@quickstart ~]$ mkdir flumeconf
[cloudera@quickstart ~]$ ls
cloudera-manager  eclipse    lib        Templates
cm_api.py         empleados.java Music      Videos
copiaempleados.java enterprise-deployment.json parcels     workspace
Desktop           express-deployment.json Pictures
Documents         flumeconf  Public
Downloads         kerberos   telefonica.java
[cloudera@quickstart ~]$ gedit flumeconf/agentefichero.log.conf
```

Fichero agentefichero.log.conf:

```
# Definición de componentes del agente
agentefichero.log.sources = source1
agentefichero.log.sinks = sink1
agentefichero.log.channels = ch1

# Configuración de propiedades del source
agentefichero.log.sources.source1.type = exec
agentefichero.log.sources.source1.command = tail -F /tmp/log.txt
agentefichero.log.sources.source1.shell = /bin/bash -c

# Configuración de propiedades del sink
agentefichero.log.sinks.sink1.type = logger

# Configuración de un canal de tipo memoria
agentefichero.log.channels.ch1.type = memory
agentefichero.log.channels.ch1.capacity = 1000
agentefichero.log.channels.ch1.transactionCapacity = 100

# Vincular source y sink al canal creado
agentefichero.log.sources.source1.channels = ch1
agentefichero.log.sinks.sink1.channel = ch1
```

```

agentefichero.log.conf
# Definición de componentes del agente
agentefichero.log.sources = source1
agentefichero.log.sinks = sink1
agentefichero.log.channels = ch1

# Configuración de propiedades del source
agentefichero.log.sources.source1.type = exec
agentefichero.log.sources.source1.command = tail -F /tmp/log.txt
agentefichero.log.sources.source1.shell = /bin/bash -c

# Configuración de propiedades del sink
agentefichero.log.sinks.sink1.type = logger

# Configuración de un canal de tipo memoria
agentefichero.log.channels.ch1.type = memory
agentefichero.log.channels.ch1.capacity = 1000
agentefichero.log.channels.ch1.transactionCapacity = 100

# Vincular source y sink al canal creado
agentefichero.log.sources.source1.channels = ch1
agentefichero.log.sinks.sink1.channel = ch1
    
```

Plain Text ▾ Tab Width: 8 ▾ Ln 8, Col 37 INS

### Preguntas:

a) ¿Qué tipo de fuente utiliza?

-Se utiliza un tipo de fuente tipo ejecutable(exe)

b) ¿Qué tipo de sumidero?

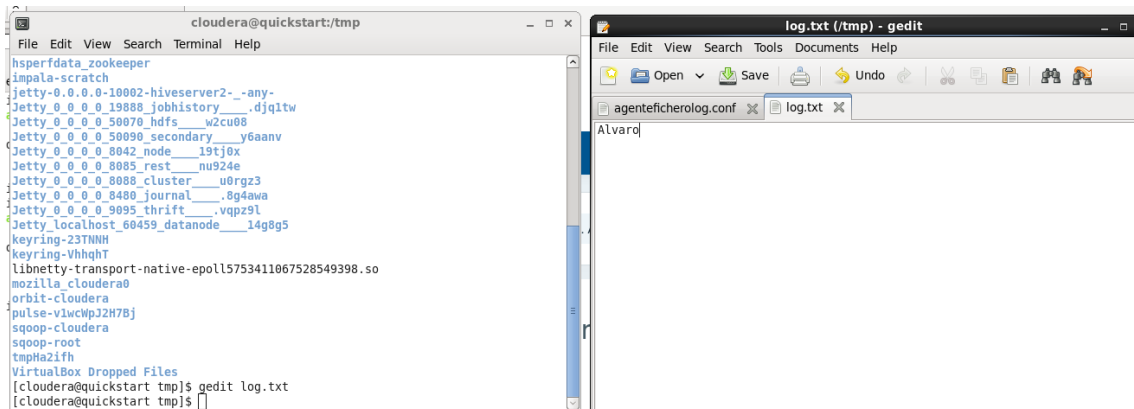
-logger

c) ¿Qué tipo de canal utiliza?

-Memoria

### Prepara el fichero log.txt:

- Crea el fichero log.txt en la ruta /tmp con el editor de texto de la máquina Cloudera.
- Añade un nombre
- Mantenlo abierto en ese editor.



### Crea el Agente Flume:

Siguiendo las indicaciones que aparecen en teoría para crear el agente Flume, construye el comando correspondiente, que tenga en cuenta lo siguiente:

- Nombre del Agente: agenteficherolog
- Ruta del Directorio de Configuración: /home/cloudera/flumeconf
- Nombre del Fichero de configuración: agenteficherolog.conf

Al lanzar el agente, debería verse en la terminal correspondiente lo siguiente:

A partir de este momento, el agente está escuchando en las operaciones que se hagan sobre el fichero /tmp/log.txt, y reaccionará a cada cambio del fichero, registrando ese evento y pintando las últimas líneas del fichero.

```
flume-ng agent - agenteficherolog -f  
/home/cloudera/flumeconf/agenteficherolog.conf  
-Dflume.root.logger=INFO,console-Xmx512m
```



## Salida:

```
[cloudera@quickstart ~]$ flume-ng agent -n agentfichero1og -f /home/cloudera/flumeconf/agentfichero1og.conf -Dflume.root.logger=INFO,console --Xmx512m
Warning: No configuration directory set! Use -conf <dir> to override.
Info: Including Hadoop libraries found via (/usr/bin/hadoop) for HDFS access
Info: Including HBASE libraries found via (/usr/bin/hbase) for HBASE access
Info: Including Hive libraries found via (/) for Hive access
+ exec (usr/java/jdk1.7.0_67-cloudera/bin/java -Xmx20m -Dflume.root.logger=INFO,console -Xmx512m -cp "/usr/lib/flume-ng/lib/*:/etc/hadoop/conf:/usr/lib/hadoop/lib/*:/usr/lib/hadoop-hdfs/*:/usr
./lib/hadoop-hdfs/lib/*:/usr/lib/hadoop-hdfs/*:/usr/lib/hadoop-yarn/lib/*:/usr/lib/hadoop-yarn/*:/usr/lib/hadoop-mapreduce/lib/*:/usr/lib/hadoop-mapreduce/*:/usr/lib/hbase/bin/*:/conf:/usr/java/jdk1.7.0_67-
cloudera/lib/tools.jar:/usr/lib/hbase/bin/*:/usr/lib/hbase/bin/*:/lib/activation-1.1.jar:/usr/lib/hbase/bin/*:/lib/apached-118n-2.0.0-M15.jar:/usr/lib/hbase/bin/*:/lib/apacheds-kerberos-codec-2.0.0-M15.jar:/usr/l
ib/hbase/bin/*:/lib/api-asn1-api-1.0.0-M20.jar:/usr/lib/hbase/bin/*:/lib/api-util-1.0.0-M20.jar:/usr/lib/hbase/bin/*:/lib/asn-3.2.jar:/usr/lib/hbase/bin/*:/lib/avro.jar:/usr/lib/hbase/bin/*:/lib/aw-aws-sdk-bundle
-1.11.134.jar:/usr/lib/hbase/bin/*:/lib/commons-beanutils-1.9.2.jar:/usr/lib/hbase/bin/*:/lib/commons-beanutils-core-1.8.0.jar:/usr/lib/hbase/bin/*:/lib/commons-cli-1.2.jar:/usr/lib/hbase/bin/*:/lib/commons-codec-1
.9.jar:/usr/lib/hbase/bin/*:/lib/commons-collections-3.2.2.jar:/usr/lib/hbase/bin/*:/lib/commons-compress-1.4.1.jar:/usr/lib/hbase/bin/*:/lib/commons-configuration-1.6.jar:/usr/lib/hbase/bin/*:/lib/commons-daemon-1
.0.13.jar:/usr/lib/hbase/bin/*:/lib/commons-digester-1.8.jar:/usr/lib/hbase/bin/*:/lib/commons-el-1.0.jar:/usr/lib/hbase/bin/*:/lib/commons-httpclient-3.1.jar:/usr/lib/hbase/bin/*:/lib/commons-io-2.4.jar:/usr/lib/h
base/bin/*:/lib/commons-lang-2.6.jar:/usr/lib/hbase/bin/*:/lib/commons-logging-1.2.jar:/usr/lib/hbase/bin/*:/lib/commons-math-2.1.jar:/usr/lib/hbase/bin/*:/lib/commons-math3-3.1.1.jar:/usr/lib/hbase/bin/*:/lib/com
ons-net-3.1.jar:/usr/lib/hbase/bin/*:/lib/core-3.1.1.jar:/usr/lib/hbase/bin/*:/lib/curator-client-2.7.1.jar:/usr/lib/hbase/bin/*:/lib/curator-framework-2.7.1.jar:/usr/lib/hbase/bin/*:/lib/curator-recipes-2.7.1.jar:/
usr/lib/hbase/bin/*:/lib/disruptor-3.3.0.jar:/usr/lib/hbase/bin/*:/lib/findbugs-annotations-1.3.9-1.jar:/usr/lib/hbase/bin/*:/lib/gson-2.2.4.jar:/usr/lib/hbase/bin/*:/lib/guava-12.0.1.jar:/usr/lib/hbase/bin/*:/lib
/hamcrest-core-1.3.jar:/usr/lib/hbase/bin/*:/lib/hbase-annotations-1.2.0-cdh5.13.0.jar:/usr/lib/hbase/bin/*:/lib/hbase-annotations-1.2.0-cdh5.13.0-tests.jar:/usr/lib/hbase/bin/*:/lib/hbase-client-1.2.0-cdh5.13.0.ja
r:/usr/lib/hbase/bin/*:/lib/hbase-common-1.2.0-cdh5.13.0.jar:/usr/lib/hbase/bin/*:/lib/hbase-common-1.2.0-cdh5.13.0-tests.jar:/usr/lib/hbase/bin/*:/lib/hbase-examples-1.2.0-cdh5.13.0.jar:/usr/lib/hbase/bin/*:/lib/h
base-external-blockcache-1.2.0-cdh5.13.0.jar:/usr/lib/hbase/bin/*:/lib/hbase-hadoop2-compat-1.2.0-cdh5.13.0.jar:/usr/lib/hbase/bin/*:/lib/hbase-hadoop2-compat-1.2.0-cdh5.13.0-tests.jar:/usr/lib/hbase/bin/*:/lib/hba
se-hadoop-compat-1.2.0-cdh5.13.0.jar:/usr/lib/hbase/bin/*:/lib/hbase-hadoop-compat-1.2.0-cdh5.13.0-tests.jar:/usr/lib/hbase/bin/*:/lib/hbase-it-1.2.0-cdh5.13.0.jar:/usr/lib/hbase/bin/*:/lib/hbase-it-1.2.0-cdh5.13.0
-tests.jar:/usr/lib/hbase/bin/*:/lib/hbase-prefix-tree-1.2.0-cdh5.13.0.jar:/usr/lib/hbase/bin/*:/lib/hbase-procedure-1.2.0-cdh5.13.0.jar:/usr/lib/hbase/bin/*:/lib/hbase-protocol-1.2.0-cdh5.13.0.jar:/usr/lib/hbase/bi
n/*:/lib/hbase-resource-bundle-1.2.0-cdh5.13.0.jar:/usr/lib/hbase/bin/*:/lib/hbase-rest-1.2.0-cdh5.13.0.jar:/usr/lib/hbase/bin/*:/lib/hbase-rsgroup-1.2.0-cdh5.13.0.jar:/usr/lib/hbase/bin/*:/lib/hbase-rsgroup-1.2.0
-cdh5.13.0-tests.jar:/usr/lib/hbase/bin/*:/lib/hbase-server-1.2.0-cdh5.13.0.jar:/usr/lib/hbase/bin/*:/lib/hbase-server-1.2.0-cdh5.13.0-tests.jar:/usr/lib/hbase/bin/*:/lib/hbase-shell-1.2.0-cdh5.13.0.jar:/usr/lib/hb
ase/bin/*:/lib/hbase-spark-1.2.0-cdh5.13.0.jar:/usr/lib/hbase/bin/*:/lib/hbase-thrift-1.2.0-cdh5.13.0.jar:/usr/lib/hbase/bin/*:/lib/high-scale-lib-1.1.1.jar:/usr/lib/hbase/bin/*:/lib/hsqldb-1.8.0.18.jar:/usr/lib/hb
ase/bin/*:/lib/htrace-core-3.2.0-incubating.jar:/usr/lib/hbase/bin/*:/lib/htrace-core4-4.0.1-incubating.jar:/usr/lib/hbase/bin/*:/lib/htrace-core.jar:/usr/lib/hbase/bin/*:/lib/httpclient-4.2.5.jar:/usr/lib/hbase/bi
```

```
File Edit View Search Terminal Help
LF43: Found binding in [jar:file:/usr/lib/flume-ng/lib/slf4j-log4j12-1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]
LF43: Found binding in [jar:file:/usr/lib/zookeeper/lib/slf4j-log4j12-1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]
LF43: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
LF43: Actual binding is of type [org.slf4j.impl.Log4jLoggerFactory]
3/04/10 11:48:24 INFO node.PollingPropertiesFileConfigurationProvider: Configuration provider starting
3/04/10 11:48:24 INFO node.PollingPropertiesFileConfigurationProvider: Reloading configuration file:/home/cloudera/flumeconf/agentfichero1og.conf
3/04/10 11:48:24 INFO conf.FlumeConfiguration: Processing:sink1
3/04/10 11:48:24 INFO conf.FlumeConfiguration: Added sinks: sink1 Agent: agentfichero1og
3/04/10 11:48:24 INFO conf.FlumeConfiguration: Processing:sink1
3/04/10 11:48:24 INFO conf.FlumeConfiguration: Post-validation flume configuration contains configuration for agents: [agentfichero1og]
3/04/10 11:48:24 INFO node.AbstractConfigurationProvider: Creating channels
3/04/10 11:48:24 INFO channel.DefaultChannelFactory: Creating instance of channel chl1 type memory
3/04/10 11:48:24 INFO node.AbstractConfigurationProvider: Created channel chl1
3/04/10 11:48:24 INFO source.DefaultSourceFactory: Creating instance of source source1, type exec
3/04/10 11:48:24 INFO sink.DefaultSinkFactory: Creating instance of sink: sink1, type: logger
3/04/10 11:48:24 INFO node.AbstractConfigurationProvider: Channel chl1 connected to [source1, sink1]
3/04/10 11:48:24 INFO node.Application: Starting new configuration: { sourceRunners: {source=org.apache.flume.source.ExecSource{name=source1,state=IDLE} } } sinkRunner
inner: { policy:org.apache.flume.sink.DefaultSinkProcessor@3a0a8d45 counterGroup: { name:null counters: { } } } channels: {chl=org.apache.flume.channel.MemoryChannel{name: chl}} }
3/04/10 11:48:24 INFO node.Application: Starting channel chl
3/04/10 11:48:24 INFO instrumentation.MonitoredCounterGroup: Monitored counter group for type: CHANNEL, name: chl: Successfully registered new MBean.
3/04/10 11:48:24 INFO node.Application: Starting Sink sink1
3/04/10 11:48:24 INFO node.Application: Starting Source source1
3/04/10 11:48:24 INFO source.ExecSource: Exec source starting with command: tail -F /tmp/log.txt
3/04/10 11:48:24 INFO instrumentation.MonitoredCounterGroup: Monitored counter group for type: SOURCE, name: source1: Successfully registered new MBean.
3/04/10 11:48:24 INFO instrumentation.MonitoredCounterGroup: Component type: SOURCE, name: source1 started
3/04/10 11:48:28 INFO sink.LoggerSink: Event: { headers:{} body: 41 6C 76 61 72 6F Alvaro }
3/04/10 11:48:28 INFO sink.LoggerSink: Event: { headers:{} body: }
3/04/10 11:48:34 INFO sink.LoggerSink: Event: { headers:{} body: 41 6C 76 61 72 6F Alvaro }
3/04/10 11:48:39 INFO sink.LoggerSink: Event: { headers:{} body: 41 6E 61 Ana }
```

